Sample document for testing the PDF LLM application.

Retrieval-Augmented Generation (RAG) combines document retrieval with language model generation.
A PDF is converted into text and split into smaller chunks. Each chunk is converted into an embedding.
The embeddings are stored in a vector database such as ChromaDB. When a user asks a question,
the application retrieves relevant chunks and sends them to a language model to create an answer.

Streamlit provides the web interface. Ollama runs language models locally.
The application should answer from the uploaded document and should say when the answer is not ava